

CS531 Programming Assignment 1: Vacuum-Cleaning Agents

Santosh Kumar, Daniel Magee
(EECS, MIME), Oregon State University

January 22, 2018

Abstract

In this project, we examine the effectiveness of three types of simple autonomous vacuum agents in two different types of environments. To compare the effectiveness of the agents, we consider the number of squares cleaned in the number of moves taken as the performance metric. We find that the random agent performed well enough to complete the course in both conditions, the simple reflex model could not complete the course, and the deterministic model could complete the open course very well but struggled on the partitioned course.

1 Introduction

We consider three types of agents in the design of an autonomous vacuum cleaner: a simple reflex agent without persistent memory, a similarly memoryless reflex agent with randomized responses to stimuli, and a model-based agent with a small amount of short term memory. Our evaluation of the three designs is based on a simple performance metric: cleaning more cells in fewer movements is preferable.

Each agent is equipped with sensors which gather three percepts from the environment: whether there is dirt in the current cell, whether there is a wall in front of the vacuum and whether the vacuum is in the cell it started in. With this information it can take one of five actions: move forward, turn off, pick up the dirt, or turn (left or right).

2 Algorithms

2.1 Memoryless Agent

The rules for the memory agent are as follows:

```
if there is DIRT then SUCK it up
if there is not DIRT and no WALL and then move FORWARD
if there is not DIRT and a WALL and I am not HOME turn RIGHT
if there is not DIRT and a WALL and I am HOME TURN OFF
```

The best possible performance for this agent is that it will pick up the dirt in every cell that abuts a wall. It may also pick up dirt on cells in the field of the room if there are posts staggered so it may run into them and turn without running along a wall. This experiment does not include this condition so, since it can only turn in corners, it only moves along walls.

2.2 Random Agent

The random agent allows for the vacuum to turn even when it is not impeded by an obstacle. A stochastic distribution is applied to the response function such that the vacuum may turn right or left based on a random choice when confronted with a wall, and may choose to turn when there is no obstacle. Moving forward when possible is generally the best decision so that choice is heavily weighted.

2.3 Memory Capable Agent

3 Results

The random agent performs better when there are no obstructions on average as it would be difficult for it to come back to the initial position and hence keeps on moving even after cleaning majority of the cells.

Fine-tuning the parameters, such as stop etc which can improve performance; as the agent stops after more or less squares are empty. Otherwise the number of moves keep on increasing if it doesnt stop and performance diminishes, probability of stopping can be increased to increase performance.

The random agent performs best when it has the space between obstacles to exercise its randomness and explore otherwise unreachable areas of

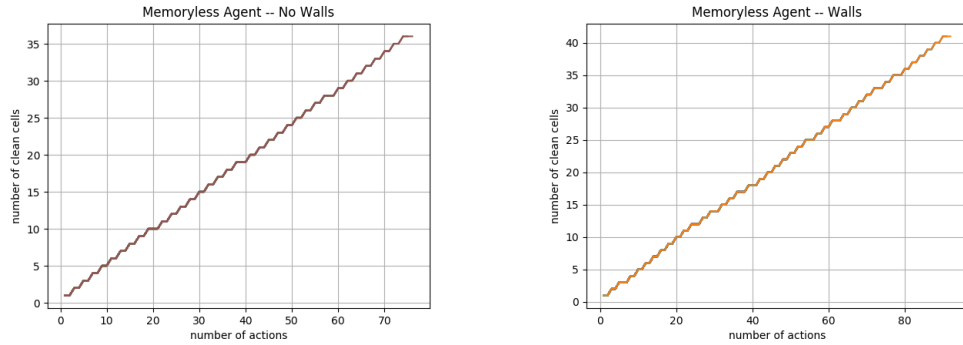


Figure 1: The Memoryless agent returns home before all the cells are clean. It runs in a loop dependent on the room configuration.

the map. More constrained areas cause a much larger deviation in the performance of the random agent which is interesting as shown in 1. While in overall performance, the random agent takes 50% longer to clean 90% of the squares on the walled map, the deviation in the time to clean that area roughly triples.




Figure 2: Averages of the results of the reflex agent with a random component

The agent with memory performs by far the best on the open map ??.




Figure 3: The agent capable of deciding movements based on recent actions

4 Conclusions

We observe several trade offs between random and deterministic agents. Deterministic agents have to store a lot of memory which is not necessary in case of randomized agents. Randomized agents do not perform well always, they only perform better on average than simple reflex agents. Randomness can be useful in cases where there are not many obstructions or else more steps taken causally without achieving desired result and the performance diminishes. Deterministic agents can be good in areas where there are lots of obstructions because storing state can be useful to avoid pitfalls in such cases, but they can also get stuck in an area if they don't have enough memory and if the target is in the starting area they may move toward it without completing the course.

If there are polygonal obstacles, we would consider implementing a mix of random and deterministic memory models where the randomness can help us deviate away from the object. The model would store such memory to repeat that path and hence steer away from the polygonal obstacle.

We were surprised that the random agent performed so well, in particular on more closed maps, and that the deterministic model performed so poorly. Using a different deterministic model, one that suited the particularities of the map would be most beneficial. We learned that more complexity does not always lead to better results, and that randomness can guide an object as well as a considered algorithm some of the time.

Table 1: Turn Random Agent cleans 90% of area.

Run	No Walls	Walls
0	483	783
1	767	755
2	457	628
3	559	641
4	667	395
5	604	1050
6	643	1115
7	499	1045
8	524	600
9	634	670
10	536	1153
11	493	677
12	465	743
13	716	522
14	680	1288
15	550	658
16	549	572
17	548	1280
18	375	472
19	620	532
20	484	720
21	554	1140
22	480	479
23	574	510
24	521	1440
25	543	525
26	423	441
27	685	686
28	404	849
29	463	1173
30	688	589
31	569	511
32	643	930
33	649	648
34	723	1046
35	598	1059
36	554	828
37	464	941
38	499	760
39	567	709
40	519	590
41	582	651
42	541	1184
43	659	771
44	566	561
45	350	753
46	330	687
47	490	789
48	522	443
49	5488	694
mean	550.02	773.72
stdev	96.01	256.95